

PolypNet-3D

Complete Project Roadmap

HACKATHON · 4 Days

PAPER · 6 Months

Easwari Engineering College · Dept. of AI & ML + Biomedical Engineering
36-Hour AI Hackathon · NIT Tiruchirappalli · Tata Centre for AI & ML

K. Jayasri · Varsha S Rao · Gokul R

Contents

PART 1 — HACKATHON	4-Day Build Plan	Page 3
Overview & Goal		3
System Architecture		3
Module 1: Frame Quality Checker		4
Module 2: YOLOv8 Detection		4
Module 3: Temporal Smoother		5
4-Day Schedule		5
Demo Strategy		6
Tech Stack		6
PART 2 — RESEARCH PAPER	6-Month Publication Plan	Page 7
Novel Contribution & Gap		7
Paper Structure		8
CRS-Net Architecture		8
New Metrics (TCS & ARS)		9
Datasets & Baselines		9
Target Journals		10
6-Month Timeline		10


PART 1 — HACKATHON BUILD

4-Day Build Plan

Build a working demo that visually proves stable, artifact-aware polyp detection. Judges want to see: Does it work? Is the problem real? Can you show the difference?

GOAL: Side-by-side demo — Plain YOLOv8 (flickering) vs PolypNet-3D (stable). That visual difference wins hackathons. No explanation needed.

System Architecture



Three modules. That's the entire system.

DM1 Frame Quality Checker — Artifact Detection

Runs BEFORE detection. Each frame gets a quality score 0–1 based on blur + brightness. Artifact frames are skipped — detection never runs on junk data.

What it catches:

Artifact Type	Signal	Detection Method
Water Washout	Frame becomes very bright (overexposed)	Brightness mean > 200
Motion Blur	Laplacian variance drops sharply	blur_score < threshold
Probe Adhesion	Dark smear across frame	Brightness mean < 40
Fecal Matter	Irregular dark patches + low contrast	Combined score < 0.3

```
def frame_quality_score(frame):  
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)  
    blur_score = cv2.Laplacian(gray, cv2.CV_64F).var() # sharpness  
    brightness = np.mean(gray) # luminance  
    blur_ok = min(blur_score / 500, 1.0)
```

```
brightness_ok = 1.0 if 40 < brightness < 200 else 0.2

return blur_ok * brightness_ok    # 0 = artifact, 1 = perfect
```

DM2 YOLOv8 Detection — Fine-tuned on KVASIR-SEG

Use pretrained YOLOv8n (nano = fastest). Fine-tune on KVASIR-SEG dataset in Google Colab T4 GPU. Takes ~2 hours overnight Day 1.

Dataset	KVASIR-SEG — 1000 annotated polyp images + videos (free, open source)
Download	datasets.simula.no/kvasir-seg
Model	YOLOv8n (nano) — fastest inference, good accuracy for demo
Training	30 epochs, batch=16, imgsz=640 — runs on free Colab T4 GPU
Time	~2 hours training + 30 min setup = done by Day 1 night

```
from ultralytics import YOLO

model = YOLO('yolov8n.pt')          # load pretrained
model.train(
    data='kvasir.yaml',              # your dataset config
    epochs=30,
    imgsz=640,
    batch=16,
    project='polypnet3d',
    name='run1'
)

# Saved to: polypnet3d/run1/weights/best.pt
```

DM3 Temporal Smoother — The Core Innovation

Three sub-components working together. This is what separates PolypNet-3D from plain YOLO.

Sub-module	What it does	Why it matters
3a. Confidence Smoother	Exponential moving average on detection confidence across frames (alpha=0.7)	Raw: 0.9→0.2→0.85→0.1 (flicker) Smoothed: 0.9→0.77→0.80→0.71 (stable)
3b. Box Holder	Holds last known bounding box for up to 8 frames during artifact frames	Polyp stays visible during washout. Doctors don't lose the detection.

3c. ByteTrack Integration	Kalman filter tracking across frames. Handles box jitter automatically.	pip install supervision 1 line of code, handles tracking fully
---------------------------	---	--

```

class BoxHolder:
    def __init__(self, hold_frames=8):
        self.last_box = None
        self.hold_counter = 0
        self.max_hold = hold_frames

    def update(self, detected_box, frame_quality):
        if frame_quality < 0.3:
            # artifact frame
            self.hold_counter += 1
            if self.hold_counter <= self.max_hold:
                return self.last_box, "HOLDING"
            # persist last box
            return None, "LOST"

        if detected_box is not None:
            self.last_box = detected_box
            self.hold_counter = 0
            return detected_box, "DETECTING"

        return None, "NO POLYP"

```

4-Day Execution Schedule

D1 Day 1 — Setup + Training

Tasks	• Install: ultralytics, supervision, opencv-python, filterpy • Download KVASIR-SEG dataset (datasets.simula.no/kvasir-seg) • Start YOLOv8n fine-tuning on Colab T4 — let run overnight • Build frame_quality_score() function and test on sample frames
Output	✓ YOLOv8 model trained (best.pt ready)

D2 Day 2 — Core Modules

Tasks	• Build Frame Quality Checker — test blur + brightness detection • Build BoxHolder class — test hold logic on sample artifact frames • Build ConfidenceSmoother — verify EMA smoothing works • Connect all three into single pipeline function
-------	---

Output	✓ Full pipeline working on test video
--------	---------------------------------------

D3 Day 3 — Demo + UI

Tasks	• Integrate ByteTrack (supervision library) for smooth box tracking • Build OpenCV overlay UI — confidence bar, quality bar, status text • Create side-by-side comparison: Plain YOLO vs PolypNet-3D • Record demo video showing flicker vs stable detection
Output	✓ Demo video recorded, side-by-side comparison ready

D4 Day 4 — Polish + Pitch

Tasks	• Fix any bugs, tune hold_frames and alpha parameters • Prepare 5 slides: Problem → Gap → Solution → Demo → Impact • Practice 3-minute pitch with demo video • Write one-paragraph abstract for submission form
Output	✓ Ready to present. Demo + slides complete.

Demo Strategy for Judges

Side-by-Side Comparison: Left screen = Plain YOLOv8 (boxes flickering, disappear during washout). Right screen = PolypNet-3D (stable box, persists through washout, smooth confidence meter).

No technical explanation needed. Judges SEE the difference in 10 seconds.

UI ELEMENT	WHAT IT SHOWS	COLOR CODE
Bounding Box	Detected polyp location	Green = high confidence, Yellow = holding
Confidence Bar	Smoothed detection confidence (EMA)	Fills left to right, 0-100%
Frame Quality Bar	Current frame cleanliness score	Green > 0.7, Red < 0.3
Status Text	DETECTING / HOLDING / ARTIFACT / LOST	Color-coded per state
Border Flash	Artifact frame detected	Red border around entire frame

Full Tech Stack

Library / Tool	Purpose	Install
ultralytics	YOLOv8 — base polyp detector	pip install ultralytics
supervision	ByteTrack + Kalman filter tracking	pip install supervision
opencv-python	Video I/O, frame processing, overlay UI	pip install opencv-python
numpy	Frame quality calculations	pip install numpy
Google Colab T4	Free GPU for model training	colab.research.google.com
KVASIR-SEG	Polyp dataset (1000 images + videos)	datasets.simula.no/kvasir-seg
SUN-SEG (optional)	Video dataset for richer testing	GitHub: GewelsJI/SUN-SEG

6-Month Publication Roadmap

Convert the hackathon demo into a publishable paper by reframing the system as a Clinical Reliability Score (CRS) framework — the first to treat polyp detection as a doctor-trust problem, not just an accuracy problem.

Paper Title: "CRS-Net: A Clinical Reliability Score Framework for Artifact-Aware Temporally-Stable Polyp Detection in Colonoscopy Video"

Target: Computers in Biology and Medicine (IF 7.0) | IEEE JBHI (IF 7.7) | Diagnostics MDPI (IF 3.0 — easiest first target)

The Research Gap — Why This Is Publishable

What Every Paper Measures	What NOBODY Has Measured (Your Contribution)
Sensitivity, Specificity, F1	Temporal Consistency Score (TCS) — your new metric
mAP on static frames	Artifact Rejection Score (ARS) — your new metric
Detection accuracy on clean frames	Detection reliability on real clinical video (artifacts included)
Model architecture novelty	Doctor trust = f(stability) — new framing of the problem
Better box on clear frames	Stable box THROUGH bad frames — no paper does this end-to-end

Clinical Evidence: A 2024 Nature study showed AI systems with 3.2% false positive rate detected adenomas in only 43.4% of cases, while a system with 0.6% FP rate reached 50.4%. Artifact frames are a primary source of false positives — nobody has built a gated system to address this.

Two New Metrics — Your Contribution to the Field

Metric	Formula	What It Measures	Why It's New
Temporal Consistency Score (TCS)	$TCS = \frac{\text{frames with consistent detection}}{\text{total frames with polyp present}}$	How stable the detection is across consecutive frames in a video	No current benchmark measures frame-to-frame consistency

Artifact Rejection Score (ARS)	$ARS = (\text{correctly skipped artifact frames}) / (\text{total artifact frames in video})$	How accurately the system identifies and skips unreliable frames	No system currently reports performance on artifact frames separately
--------------------------------	--	--	---

Full Paper Structure

1. Introduction (1–1.5 pages)	- Colorectal cancer = 2nd leading cancer death globally - AI detection systems exist but doctors don't fully trust them - Why? False positives from artifacts + flickering detections - Gap: No system measures or optimizes for clinical reliability - Contributions: CRS-Net architecture + TCS metric + ARS metric
2. Related Work (1.5 pages)	- Frame-level detection: YOLO variants, anchor-free detectors - Temporal detection: TSdetector (2025), 3D-FCN, Hybrid 2D/3D CNN - Artifact handling: None exist end-to-end — your gap - Table: Prior works vs yours across all dimensions
3. Methodology — CRS-Net (2 pages)	3a. Clinical Reliability Score: weighted combination of blur + brightness + motion (weights learned from data, not fixed — this is deeper than hackathon) 3b. Gated Detection: if CRS > threshold → detect; else → temporal interpolation 3c. Temporal Stability Module: Kalman filter + confidence EMA + re-ID after artifact 3d. Architecture diagram showing full flow with formulas
4. Experiments (2 pages)	- Datasets: KVASIR-SEG + SUN-SEG (49,136 frames) + CVC-ClinicDB - Baselines: Plain YOLOv8 TSdetector 3D-FCN Your CRS-Net - Metrics: F1, mAP, Sensitivity, Specificity + TCS + ARS + FP per video - Ablation study: each component added one by one showing improvement
5. Results & Discussion (1.5 pages)	- Quantitative table: CRS-Net vs all baselines on all metrics - Qualitative figures: frame-by-frame during washout (visual comparison) - Clinical meaning: what does 10% better TCS mean for endoscopists? - Limitation: trained on limited artifact data — future work needed
6. Conclusion (0.5 pages)	- First system to frame polyp detection as a clinical trust problem - Two new metrics introduced for the field - Future: extend to other endoscopic procedures (gastroscopy, capsule)

CRS-Net: Core Formula

```
# Clinical Reliability Score (CRS) – learned weights, not fixed
CRS(t) = w1 * BlurScore(t) + w2 * BrightnessScore(t) + w3 * MotionScore(t)

# Gated Detection
if CRS(t) > threshold:
    detection = YOLO(frame_t)          # normal detection
else:
    detection = interpolate(t-1, t+1)  # temporal interpolation

# Temporal Stability
smooth_conf(t) = alpha * smooth_conf(t-1) + (1 - alpha) * raw_conf(t)

# Final Output per frame
output = {
    'box':      kalman_smooth(detection.box),
    'confidence': smooth_conf(t),
    'crs':      CRS(t),
    'status':   'DETECTING' | 'HOLDING' | 'ARTIFACT' | 'LOST'
}
```

Expected Ablation Study Results

Model Configuration	F1 Score	TCS	ARS	FP / Video
Baseline YOLOv8 (no temporal)	0.84	0.61	—	4.2
+ Temporal smoothing only	0.86	0.78	—	3.8
+ CRS gating only	0.87	0.74	0.79	2.1
+ Confidence EMA only	0.85	0.82	—	3.5
CRS-Net (all components)	0.89	0.91	0.88	1.3
TSdetector (2025 SOTA)	0.88	0.83	—	2.9

★ CRS-Net beats TSdetector (2025 SOTA) on all metrics — that's your publishable claim.

Datasets & Baselines

Dataset	Size	Type	Use
KVASIR-SEG	1,000 images	Static images	Primary training + fine-tuning

SUN-SEG	49,136 frames / 285 clips	Video colonoscopy	Temporal evaluation (TCS, ARS)
CVC-ClinicDB	612 images	Static images	Benchmark comparison
SUN-SEG Negative	~1,000 artifact clips	Video with artifacts	ARS evaluation specifically

Target Journals

Journal	Impact Factor	Scope Match	Difficulty	Recommended?
Diagnostics (MDPI)	3.0	High	★■	Start here — fastest
Computers in Biology and Medicine (Elsevier)	7.0	Very High	★★■	Primary target
IEEE JBHI	7.7	High	★★★	After 1 revision
Medical Image Analysis	10.7	High	★★★★	Top tier — long-term

6-Month Paper Timeline

AUG — Hackathon → Research Code	
Tasks	Convert demo to clean research codebase. Add learned CRS weights. Set up experiment framework for reproducibility.
Output	✓ Clean GitHub repo with training scripts
SEP — Experiments on SUN-SEG	
Tasks	Run full evaluation on SUN-SEG (49K frames). Compare against TSdetector and baselines. Compute TCS and ARS for all models. Identify artifact frame categories.
Output	✓ All experiment results logged
OCT — Ablation + Analysis	
Tasks	Run ablation study — each component added step by step. Generate qualitative figures (side-by-side frames during washout). Statistical significance testing.
Output	✓ All tables and figures ready

NOV

NOV — Write Paper

Tasks	Write all 6 sections. Introduction → Related Work → Methodology → Experiments → Results → Conclusion. Get feedback from Dr. G. Babu.
Output	✓ Full paper draft complete

DEC

DEC — Submit

Tasks	Final proofreading, format for target journal. Submit to Computers in Biology and Medicine or Diagnostics. Aligns with your placement readiness deadline.
Output	✓ Submitted ✓ — Placement portfolio complete

Quick Comparison — Hackathon vs Paper

Dimension	Hackathon (4 Days)	Research Paper (6 Months)
Goal	Win prize with working demo	Publish in indexed journal
CRS Weights	Fixed manually (simple)	Learned from data (rigorous)
Datasets	KVASIR-SEG only	KVASIR + SUN-SEG + CVC-ClinicDB
Metrics	Visual demo comparison	F1, mAP, TCS, ARS, FP/video
Novelty Claim	Artifact rejection demo	New CRS framework + 2 new metrics
Audience	Hackathon judges (5 min)	International peer reviewers
Time	4 days	4–6 months
Output	Demo video + 5 slides	8-page conference/journal paper
Who Cares?	Judges want visual impact	Reviewers want quantitative novelty

One-Line Hackathon Pitch:

"Unlike TSdetector which still fails on artifact frames, PolypNet-3D introduces artifact-aware temporal holding — the first system that maintains stable clinical overlay even during water washout and probe adhesion frames."

One-Line Paper Pitch:

"We propose CRS-Net, a clinical reliability score framework that gates detection on artifact frames and introduces two new metrics — TCS and ARS — enabling the first quantitative measurement of doctor-trust reliability in AI-assisted colonoscopy."

Same system. Two different stories. Build the hackathon version first — then scale it into the paper after December.